

PS 312: Programming with R

Course Notes

Abhijit Dasgupta, PhD

Last updated: March 27, 2019

Welcome

This course is an introduction to the statistical programming language [R](#) and various applications. We will cover the entire data analytics pipeline from data ingestion to data wrangling, summarizing, modeling, visualizing and reporting, all using tools found within the R ecosystem.

The version of these notes you are reading now was built on 2019-03-27.

Reproducibility

These notes are written with [bookdown](#), a R package for writing books using [rmarkdown](#). All code in these notes were developed on R version 3.5.0 (2018-04-23), using the same packages pre-installed in your virtual machines. When you're on your own, you will need to install a recent version of R, and also install the corresponding packages, on your computer, for all the code to work. A listing of all the packages used in this course will be available as an appendix.

To build these notes locally, clone or [download](#) the [Github repo](#) hosting these notes, unzip it if necessary, and double-click on `FSI_Book.Rproj`. Assuming you have RStudio installed, this will open this project (more on *RStudio Projects* later). You can then go to the console and enter the following code:

```
bookdown::render_book("index.Rmd") # to build these notes
browseURL("_book/index.html") # to view it
```

Chapter 1

Data visualization

1.1 ggplot2

We'll be primarily using ggplot2 in this workshop.

- Makes pretty good formatting choices out of the box
- Works like pipes!!
- Is declarative (tell it what you want) without getting caught up in minutiae
- Strongly leverages data frames (good practice)
- Fast enough
- There are good templates if you want to change the look

The ggplot2 package is a very flexible and (to me) intuitive way of visualizing data. It is based on the concept of layering elements on a canvas.


This idea of layering graphics on a canvas is, to me, a nice way of building graphs

- A `data.frame` object
- *Aesthetic mappings* (`aes`) to say what data is used for what purpose in the viz
 - x- and y-direction
 - shapes, colors, lines
- A *geometry object* (`geom`) to say what to draw
 - You can “layer” geoms on each other to build plots

ggplot used pipes before pipes were a thing.


However, it uses the `+` symbol for piping rather than the `%>%` operator, since it pre-dates the tidyverse

```
library(ggplot2)
ggplot(mtcars, aes(x = wt, y = mpg)) + geom_point()
```



- A `data.frame` object: `mtcars`
- Aesthetic mapping:
 - x-axis: `wt`
 - y-axis: `mpg`
- Geometry:
 - `geom_point`: draw points


```
ggplot(mtcars, aes(x = wt, y = mpg)) +  
  geom_point() +  
  geom_smooth()
```



- A `data.frame` object: `mtcars`
- Aesthetic mapping:
 - x-axis: `wt`
 - y-axis: `mpg`
- Geometry:
 - `geom_point`: draw points
 - `geom_smooth`: Add a layer which draws a best-fitting line

Now we clean up the plot to make it presentable.

Fuel efficiency by weight

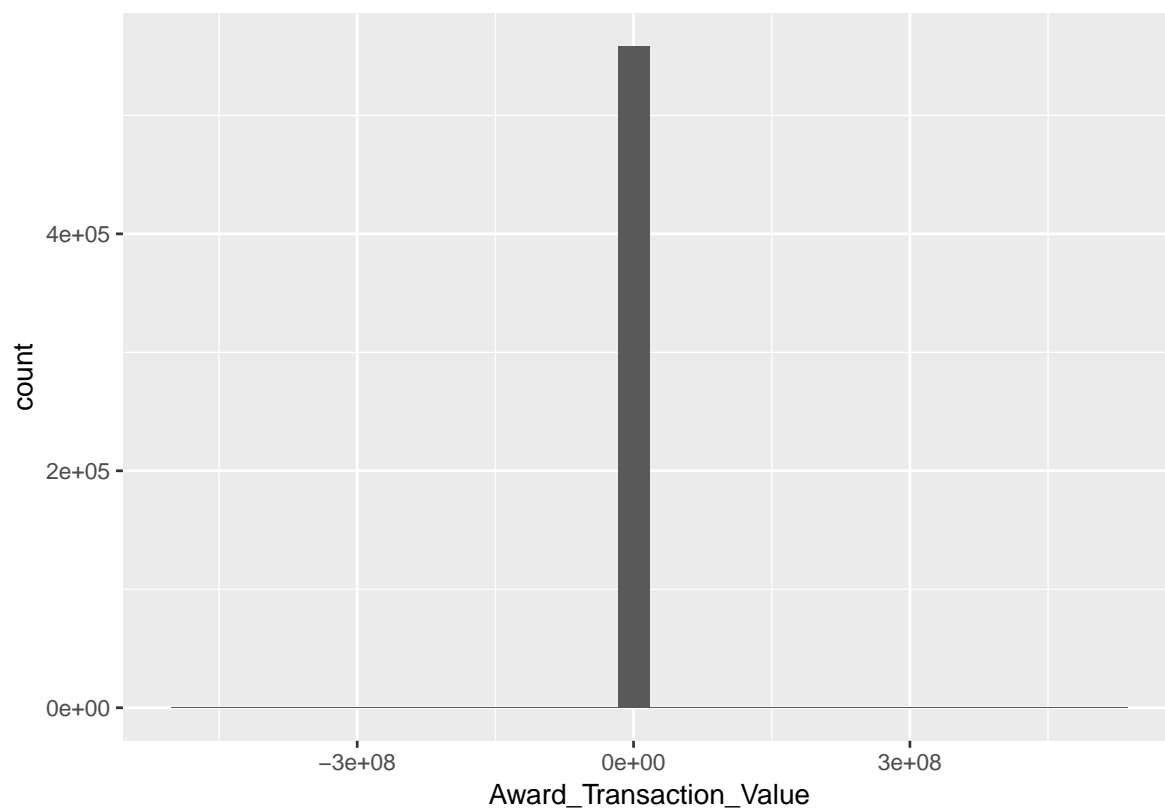


Source: mtcars dataset

1.2 Single continuous variable

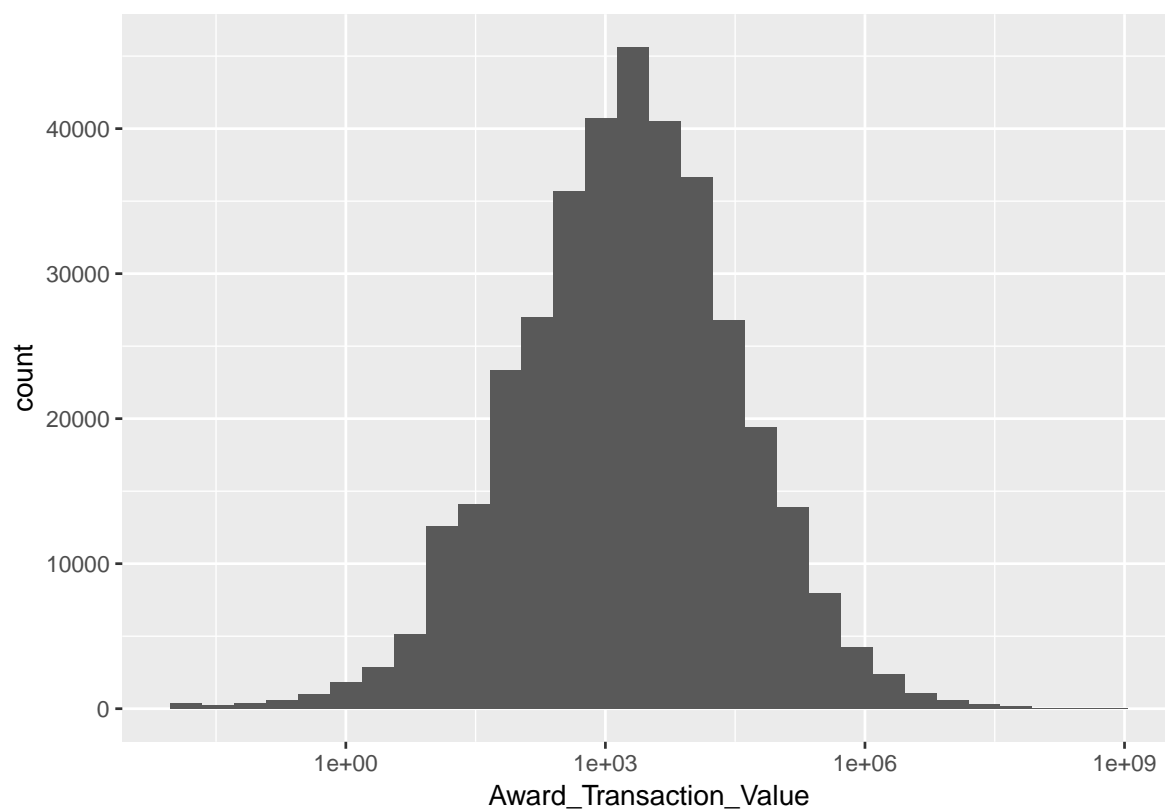
1.2.1 Histogram

```
dos <- import('data/Department of State.csv')
dos %>%
  ggplot(aes(x = Award_Transaction_Value)) + geom_histogram()
```



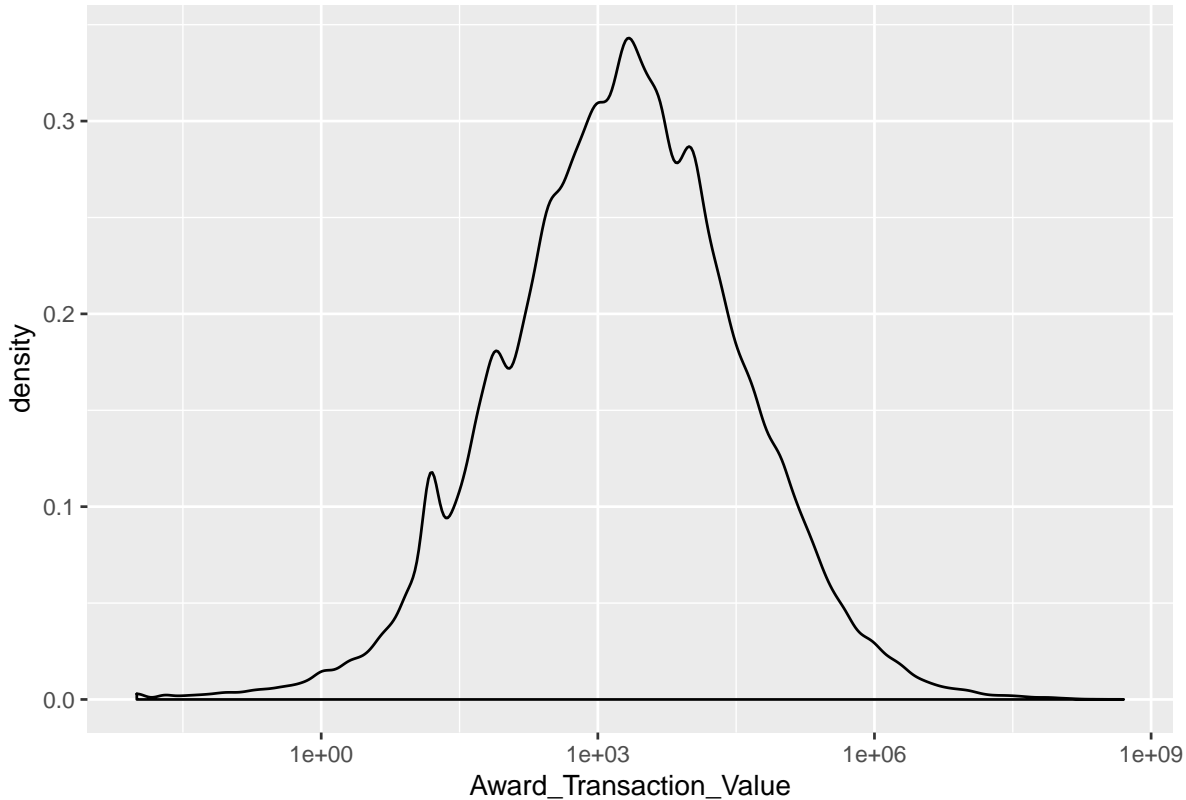
Change the axis to the log scale for better visual

```
dos %>%  
  ggplot(aes(x = Award_Transaction_Value)) + geom_histogram() +  
  scale_x_log10() # x-axis on log scale
```




1.2.2 Density plot

```
dos %>%  
  ggplot(aes(x = Award_Transaction_Value)) + geom_density() +  
  scale_x_log10()
```



1.3 Bar plots

```
library(lubridate)  
dos %>%  
  group_by(year = year(as_date(Award_Start_Date))) %>%  
  summarize(amount = sum(Award_Transaction_Value)) %>%  
  ggplot(aes(x = year, y = amount)) + # Note change in pipe operator  
  geom_bar(stat='identity')
```





1.3.1 Exercise

Using the `mtcars` dataset in R, create:


1. A histogram of the fuel efficiencies (`mpg`) in the data set
2. A bar plot of frequencies of number of cylinders (`cyl`) in the car

```
ggplot(mtcars, aes(x = mpg)) + geom_histogram(binwidth=3)
```



```
# ggplot(mtcars) + geom_histogram(aes(x = mpg), binwidth = 3)
```


```
ggplot(mtcars, aes(x = factor(cyl))) + geom_bar()
```



1.4 Two continuous variables


1.4.1 Adding a best fitting straight line

```
ggplot(mtcars, aes(x = hp, y = mpg)) +  
  geom_point() +  
  geom_smooth(method = 'lm')
```



1.5 Time series

```
library(forecast)  
d <- data.frame(x = 1:length(gas), y = gas) # Australian monthly gas production  
ggplot(d, aes(x, y)) + geom_line()
```




1.5.1 Exercise

1. Create a scatter plot of sepal length and sepal width from the `iris` dataset, and add a smooth line through it

1.6 Continuous variable with discrete variable

1.6.1 Boxplot


```
dos %>%
  ggplot(aes(x = factor(year(as_date(Award_Start_Date))),
             y = Award_Transaction_Value)) +
  geom_boxplot() +
  scale_y_log10() +
  labs(x = 'Year')
```



1.6.2 Violin plot

This is essentially a reflected density plot and gives a better sense of the data distribution

```
dos %>%
  ggplot(aes(x = factor(year(as_date(Award_Start_Date))),
              y = Award_Transaction_Value)) +
  geom_violin() +
  scale_y_log10() +
  labs(x = 'Year')
```




1.6.3 Exercise

1. Plot a boxplot of petal length by species using the `iris` dataset


1.7 Grouped visualizations

We're going to plot the change in aid provided to each country over time. To do this we need summaries by time and location

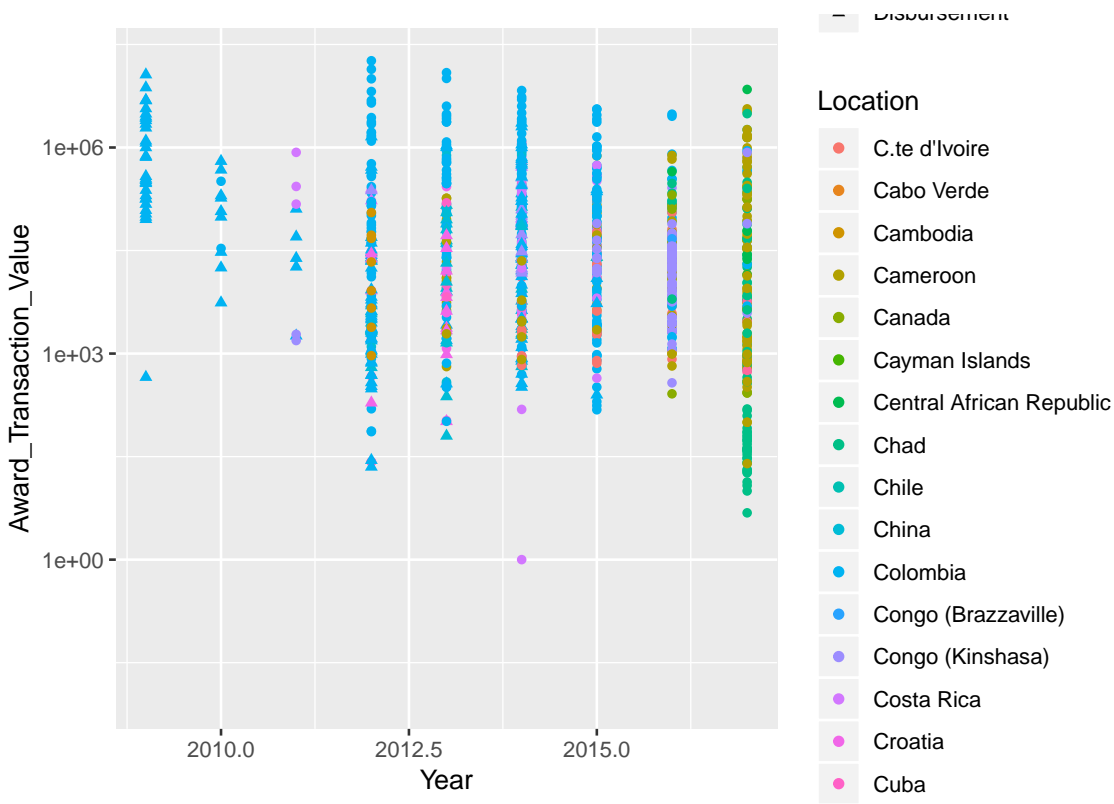
```
grp_data <- dos %>%
  group_by(Recipient_Location, year = year(as_date(Award_Start_Date))) %>%
  summarize(amt = sum(Award_Transaction_Value)) %>%
  filter(str_detect(Recipient_Location, '^C'))
ggplot(grp_data, aes(x = year, y = amt, color=Recipient_Location))+
  geom_line()+
  scale_y_log10()
```




```
ggplot(grp_data, aes(x = year, y = amt)) +
  geom_line() +
  scale_y_log10() +
  facet_wrap(~Recipient_Location)
```



```
## dos %>% filter(str_detect(Recipient_Location, '^C')) %>%
## ggplot(aes(x = year(as_date(Award_Start_Date)),
##           y = Award_Transaction_Value,
##           color = Recipient_Location,
##           shape = Award_Transaction_Type))+
##   geom_point()+
##   labs(x = 'Year', color='Location')+
##   scale_y_log10()
##
```



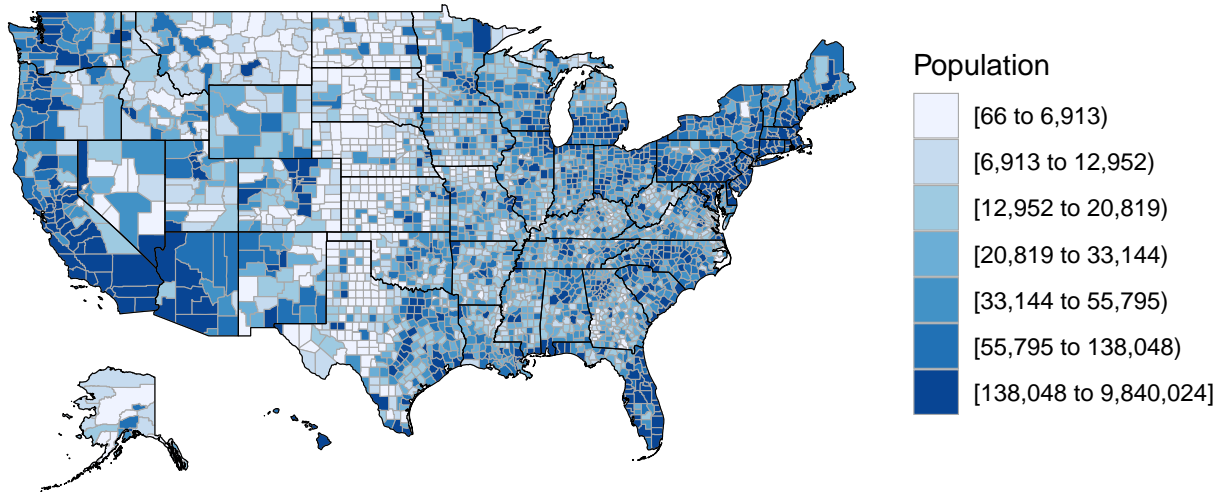
```
## dos %>% filter(str_detect(Recipient_Location, '^C')) %>%
## ggplot(aes(x = year(as_date(Award_Start_Date)),
##           y = Award_Transaction_Value,
##           color = Recipient_Location,
##           shape = Award_Transaction_Type))+
##   geom_jitter()+
##   labs(x = 'Year', color='Location')+
##   scale_y_log10()
##
```

```
schools <- rio::import('data/schools.rds')
schools %>% filter(tophead=='Elementary schools',
                  head2=="Average hours in school day") %>%
  filter(!is.na(State), State != 'United States') %>%
  ggplot(aes(x = State, y = stats, ymin = stats - 2*se,
             ymax = stats + 2*se)) +
  geom_pointrange() +
  labs(y = 'Avg hours in school day') +
  theme_bw() +
  theme(axis.text.x = element_text(angle=45, hjust = 1))
```

1.8 Maps

US 2012 County Population Estimates



We can also ingest SHP files to draw maps. We don't show the final version since it took too long to render.

```
library(sf)
hrr_info <- st_read('~Downloads/hrr_bdry-1/HRR_Bdry.SHP')
head(hrr_info)
ggplot(hrr_info) + geom_sf()
ggsave('map.png')
```

1.9 Stitching graphs together.

```
# install.packages('cowplot')
library(cowplot)
p1 <- ggplot(iris, aes(Sepal.Length, Sepal.Width, color = Species)) +
  geom_point() + facet_grid(. ~ Species) + stat_smooth(method = "lm") +
  background_grid(major = 'y', minor = "none") +
  panel_border() + theme(legend.position = "none")

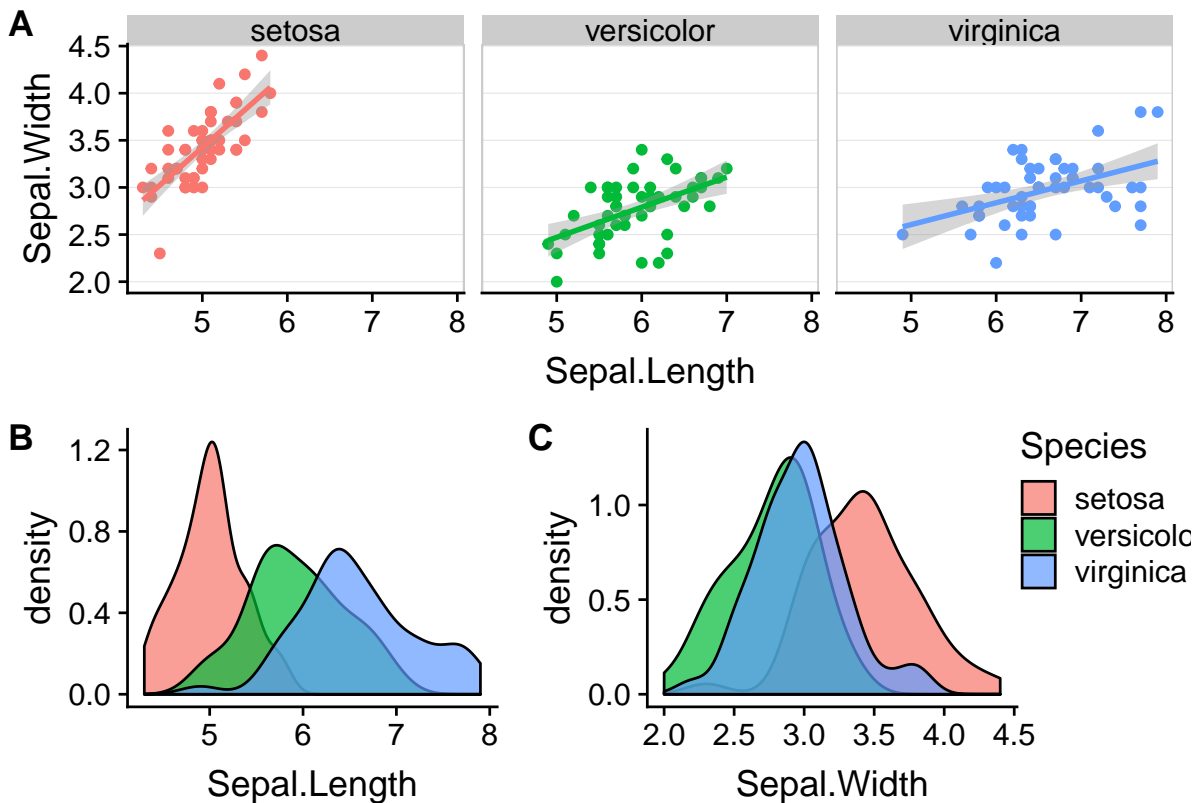
# plot B
p2 <- ggplot(iris, aes(Sepal.Length, fill = Species)) +
  geom_density(alpha = .7) + theme(legend.justification = "top")
p2a <- p2 + theme(legend.position = "none")

# plot C
p3 <- ggplot(iris, aes(Sepal.Width, fill = Species)) +
  geom_density(alpha = .7) + theme(legend.position = "none")

# legend
legend <- get_legend(p2)
```

```
# align all plots vertically
plots <- align_plots(p1, p2a, p3, align = 'v', axis = 'l')

# put together bottom row and then everything
bottom_row <- plot_grid(plots[[2]], plots[[3]], legend, labels = c("B", "C"), rel_widths = c(1, 1))
plot_grid(plots[[1]], bottom_row, labels = c("A"), ncol = 1)
```



```
## library(ggplot2)
## library(plotly)
## p=ggplot(iris, aes(x=Sepal.Length,
##                    y=Sepal.Width,
##                    color=Species,
##                    shape=Species)) +
##   geom_point(size=6, alpha=0.6)
## mytext=paste("Sepal Length = ", iris$Sepal.Length,
##              "\n", "Sepal Width = ", iris$Sepal.Width,
##              "\n", "Row Number: ", rownames(iris), sep="")
## pp=plotly::plotly_build(p)
## style( pp, text=mytext,
##        hoverinfo = "text",
##        traces = c(1, 2, 3) )
```

1.10 Interactive graphics

We won't put these in the notes, since they don't work well in printed form

Chapter 2

Functions

```
myDumbFunction <- function() 42  
myDumbFunction()
```

```
[1] 42
```

```
doubleIt <- function(x) {  
  myResult <- x * 2  
  myResult # or, explicitly, return(myResult)  
}  
doubleIt(5)
```

```
[1] 10
```

```
exists("myResult")
```

```
[1] FALSE
```

```
myResult <- 1000  
doubleItOutput <- doubleIt(2)  
myResult
```

```
[1] 1000
```

```
my_sum <- function(x){  
  s <- sum(x)  
  n <- length(x)  
  result <- s / n  
  return(result)  
}
```

```
my_sum(1:10)
```

```
[1] 5.5
```

```
answer <- my_sum(1:10)  
answer
```

```
[1] 5.5
```

```
my_sum <- function(x){
  s <- sum(x)
  n <- length(x)
  results<- list(sum = s, length = n, answer = s / n)
  return(results)
}
```

```
my_sum(1:10)
```

```
$sum
[1] 55
```

```
$length
[1] 10
```

```
$answer
[1] 5.5
```

```
my_sum <- function(x){
  s <- sum(x)
  n <- length(x)
  results<- list(sum = s, length = n, answer = s / n)
  return(results)
}
```

```
answer <- my_sum(1:10)
answer$answer
```

```
[1] 5.5
```

```
answer[['answer']]
```

```
[1] 5.5
```

```
names(answer)
```

```
[1] "sum"      "length" "answer"
```

```
x <- 1:10
x[3] <- NA
my_sum(x)
```

```
$sum
[1] NA
```

```
$length
[1] 10
```

```
$answer
[1] NA
```

```
my_sum <- function(x){
  s <- sum(x, na.rm=T)
  n <- length(!is.na(x))
  results <- list("sum" = s, "length" = n, "answer" = s/n)
}
my_sum(x)
```

```
my_sum <- function(x){
  s <- sum(x, na.rm = T)
  n <- length(!is.na(x))
  results <- list("sum" = s, "length" = n, "answer" = s/n)
  return(results) #<<
}
my_sum(x)
```

```
$sum
[1] 52
```

```
$length
[1] 10
```

```
$answer
[1] 5.2
```

```
my_sum <- function(x){
  s <- sum(x, na.rm = T)
  n <- length(!is.na(x))
  results <- list("sum" = s, "length" = n, "answer" = s/n)
  return(results)
}
```

```
my_sum <- function(x){
  s <- sum(x, na.rm = T)
  {{ n <- sum(!is.na(x)) }}
  results <- list("sum" = s, "length" = n, "answer" = s/n)
  return(results)
}
my_sum(x)
```

```
$sum
[1] 52
```

```
$length
[1] 9
```

```
$answer
[1] 5.777778
```

```
my_sum <- function(x){
  s <- sum(x, na.rm = T)
  n <- sum(!is.na(x))
  results <- list("sum" = s, "length" = n, "answer" = s/n)
  return(results)
}
```

```
my_sum <- function(x, remove_missing = TRUE){ #<<
  s <- sum(x, na.rm = T)
  n <- sum(!is.na(x))
  results <- list("sum" = s, "length" = n, "answer" = s/n)
  return(results)
}
```

```
my_sum <- function(x, remove_missing = TRUE){
  {{if(remove_missing){
    x <- x[!is.na(x)]
  }}
  s <- sum(x)
  n <- length(x)}}
  results <- list("sum" = s, "length" = n, "answer" = s/n)
  return(results)
}
my_sum(x)
```

```
$sum
[1] 52
```

```
$length
[1] 9
```

```
$answer
[1] 5.777778
```

```
my_sum <- function(x, remove_missing = TRUE){
  if(remove_missing){
    x <- x[!is.na(x)]
  }
  s <- sum(x)
  n <- length(x)
  results <- list("sum" = s, "length" = n, "answer" = s/n, "nmiss" = sum(is.na(x)))
  return(results)
}
my_sum(x)
```

```
$sum
[1] 52
```



```
$length
[1] 9
```

```
$answer
[1] 5.777778
```

```
$nmiss
[1] 0
```

```
my_sum <- function(x, remove_missing = TRUE){
  nmiss <- sum(is.na(x)) #<<
  if(remove_missing){
    x <- x[!is.na(x)]
  }
  s <- sum(x)
  n <- length(x)
  results <- list("sum" = s, "length" = n, "answer" = s/n, "nmiss" = sum(is.na(x)))
  return(results)
}
my_sum(x)
```

```
$sum
[1] 52
```

```
$length
[1] 9
```

```
$answer
[1] 5.777778
```

```
$nmiss
[1] 0
```

```
my_sum <- function(x, remove_missing = TRUE){
  nmiss <- sum(is.na(x))
  if(remove_missing){
    x <- x[!is.na(x)]
  }
  s <- sum(x)
  n <- length(x)
  results <- list("sum" = s, "length" = n, "answer" = s/n, "nmiss" = nmiss) #<<
  return(results)
}
my_sum(x)
```

```
$sum
[1] 52
```

```
$length
```

```
[1] 9
```

```
$answer
```

```
[1] 5.777778
```

```
$nmiss
```

```
[1] 1
```

```
my_sum(x, remove_missing = F)
```

```
$sum
```

```
[1] NA
```

```
$length
```

```
[1] 10
```

```
$answer
```

```
[1] NA
```

```
$nmiss
```

```
[1] 1
```

```
my_summary <- function(d){
```

```
}
```

```
my_summary <- function(d){
```

```
  require(tidyverse) #<
```

```
}
```

```
my_summary <- function(d){
```

```
  require(tidyverse)
```

```
  summary_cts <- d %>%
```

```
    summarize_if(is.numeric, list("mean" = ~mean(x, na.rm=T),
                                   "median" = ~median(x, na.rm=T),
                                   'sd' = ~sd(x, na.rm=T),
                                   'nmiss' = ~sum(is.na(x))))
```

```
  return(list("cts" = summary_cts))
```

```
}
```

```
my_summary(iris)
```

```
Loading required package: tidyverse
```

```
-- Attaching packages ----- tidyverse 1.2.1 --
```

```
v ggplot2 3.1.0      v purrr   0.3.2
v tibble   2.0.1      v dplyr   0.8.0.9009
v tidyr    0.8.3      v stringr 1.4.0
v readr    1.3.1      v forcats 0.4.0
```

```
Warning: package 'tibble' was built under R version 3.5.2
```

Warning: package 'tidyr' was built under R version 3.5.2

Warning: package 'stringr' was built under R version 3.5.2

-- Conflicts ----- tidyverse_conflicts() --

x dplyr::filter() masks stats::filter()

x dplyr::lag() masks stats::lag()

\$cts

	Sepal.Length_mean	Sepal.Width_mean	Petal.Length_mean	Petal.Width_mean
1	5.777778	5.777778	5.777778	5.777778
	Sepal.Length_median	Sepal.Width_median	Petal.Length_median	
1	6	6	6	
	Petal.Width_median	Sepal.Length_sd	Sepal.Width_sd	Petal.Length_sd
1	6	3.073181	3.073181	3.073181
	Petal.Width_sd	Sepal.Length_nmiss	Sepal.Width_nmiss	Petal.Length_nmiss
1	3.073181	1	1	1
	Petal.Width_nmiss			
1	1			

```
my_summary <- function(d){
  require(tidyverse)
  summary_cts <- d %>%
    summarize_if(is.numeric, list("mean" = ~mean(x, na.rm=T),
                                   "median" = ~median(x, na.rm=T),
                                   'sd' = ~sd(x, na.rm=T),
                                   'nmiss' = ~sum(is.na(x)))) %>%
    gather(variable, value) %>%
    separate(variable, c("variable", "stat"), sep='_') %>%
    spread(stat, value)
  return(list("cts" = summary_cts))
}
my_summary(iris)
```

\$cts

	variable	mean	median	nmiss	sd
1	Petal.Length	5.777778	6	1	3.073181
2	Petal.Width	5.777778	6	1	3.073181
3	Sepal.Length	5.777778	6	1	3.073181
4	Sepal.Width	5.777778	6	1	3.073181

```
my_summary <- function(d){
  require(tidyverse)
  summary_cts <- d %>%
    summarize_if(is.numeric, list("mean" = ~mean(x, na.rm=T), #<<
                                   "median" = ~median(x, na.rm=T), #<<
                                   'sd' = ~sd(x, na.rm=T), #<<
                                   'nmiss' = ~sum(is.na(x)))) %>% #<<
    gather(variable, value) %>%
    separate(variable, c("variable", "stat"), sep='_') %>%
```

```

    spread(stat, value)
  return(list("cts" = summary_cts))
}
my_summary(iris)

```

```
$cts
```

	variable	mean	median	nmiss	sd
1	Petal.Length	5.777778	6	1	3.073181
2	Petal.Width	5.777778	6	1	3.073181
3	Sepal.Length	5.777778	6	1	3.073181
4	Sepal.Width	5.777778	6	1	3.073181

```

my_summary <- function(d){
  require(tidyverse)
  summary_cts <- d %>%
    summarize_if(is.numeric, list("mean" = ~mean(., na.rm=T), #<<
                                   "median" = ~median(., na.rm=T), #<<
                                   'sd' = ~sd(., na.rm=T), #<<
                                   'nmiss' = ~sum(is.na(.))) %>% #<<

    gather(variable, value) %>%
    separate(variable, c("variable", "stat"), sep='_') %>%
    spread(stat, value)
  return(list("cts" = summary_cts))
}
my_summary(iris)

```

```
$cts
```

	variable	mean	median	nmiss	sd
1	Petal.Length	3.758000	4.35	0	1.7652982
2	Petal.Width	1.199333	1.30	0	0.7622377
3	Sepal.Length	5.843333	5.80	0	0.8280661
4	Sepal.Width	3.057333	3.00	0	0.4358663

```

my_summary <- function(d){
  require(tidyverse)
  summary_cts <- d %>%
    summarize_if(is.numeric, list("mean" = ~mean(., na.rm=T),
                                   "median" = ~median(., na.rm=T),
                                   'sd' = ~sd(., na.rm=T),
                                   'nmiss' = ~sum(is.na(.))) %>%

    gather(variable, value) %>%
    separate(variable, c("variable", "stat"), sep='_') %>%
    spread(stat, value) %>%
    select(variable, nmiss, everything()) #<<
  return(list("cts" = summary_cts))
}
my_summary(iris)

```



```

                                'nmiss' = ~sum(is.na(.))) %>%
gather(variable, value) %>%
separate(variable, c("variable", "stat"), sep='_') %>%
spread(stat, value) %>%
select(variable, nmiss, everything())
summary_cat <- d %>%
  summarise_if(is.factor, list('nmiss' = ~sum(is.na(.)),
                                'ncat' = ~length(unique(.)),
                                'categories' = ~paste(sort(unique(levels(.))), collapse=","))
  )
return(list("cts" = summary_cts,
            "cat" = summary_cat))
}
my_summary(x)

```

```

datas <- list('cars' = mtcars, 'iris' = iris, 'diamonds' = diamonds)
map(datas, my_summary)

```

\$cars

\$cars\$cts

	variable	nmiss	mean	median	sd
1	am	0	0.406250	0.000	0.4989909
2	carb	0	2.812500	2.000	1.6152000
3	cyl	0	6.187500	6.000	1.7859216
4	disp	0	230.721875	196.300	123.9386938
5	drat	0	3.596563	3.695	0.5346787
6	gear	0	3.687500	4.000	0.7378041
7	hp	0	146.687500	123.000	68.5628685
8	mpg	0	20.090625	19.200	6.0269481
9	qsec	0	17.848750	17.710	1.7869432
10	vs	0	0.437500	0.000	0.5040161
11	wt	0	3.217250	3.325	0.9784574

\$cars\$cat

data frame with 0 columns and 1 row

\$iris

\$iris\$cts

	variable	nmiss	mean	median	sd
1	Petal.Length	0	3.758000	4.35	1.7652982
2	Petal.Width	0	1.199333	1.30	0.7622377
3	Sepal.Length	0	5.843333	5.80	0.8280661
4	Sepal.Width	0	3.057333	3.00	0.4358663

\$iris\$cat

nmiss	ncat	categories

```
1      0      3 setosa, versicolor, virginica
```

```
$diamonds
```

```
$diamonds$cts
```

```
# A tibble: 7 x 5
```

	variable	nmiss	mean	median	sd
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	carat	0	0.798	0.7	0.474
2	depth	0	61.7	61.8	1.43
3	price	0	3933.	2401	3989.
4	table	0	57.5	57	2.23
5	x	0	5.73	5.7	1.12
6	y	0	5.73	5.71	1.14
7	z	0	3.54	3.53	0.706

```
$diamonds$cat
```

```
# A tibble: 1 x 9
```

	cut_nmiss	color_nmiss	clarity_nmiss	cut_ncat	color_ncat	clarity_ncat
	<int>	<int>	<int>	<int>	<int>	<int>
1	0	0	0	5	7	8

```
# ... with 3 more variables: cut_categories <chr>, color_categories <chr>,
#   clarity_categories <chr>
```

Chapter 3

Modeling

```
library(survival)
```

```
data(pbc)
```

```
str(pbc)
```

```
'data.frame':  418 obs. of  20 variables:
 $ id      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ time    : int  400 4500 1012 1925 1504 2503 1832 2466 2400 51 ...
 $ status  : int  2 0 2 2 1 2 0 2 2 2 ...
 $ trt     : int  1 1 1 1 2 2 2 2 1 2 ...
 $ age     : num  58.8 56.4 70.1 54.7 38.1 ...
 $ sex     : Factor w/ 2 levels "m","f": 2 2 1 2 2 2 2 2 2 2 ...
 $ ascites : int  1 0 0 0 0 0 0 0 0 1 ...
 $ hepato  : int  1 1 0 1 1 1 1 0 0 0 ...
 $ spiders : int  1 1 0 1 1 0 0 0 1 1 ...
 $ edema   : num  1 0 0.5 0.5 0 0 0 0 0 1 ...
 $ bili    : num  14.5 1.1 1.4 1.8 3.4 0.8 1 0.3 3.2 12.6 ...
 $ chol    : int  261 302 176 244 279 248 322 280 562 200 ...
 $ albumin : num  2.6 4.14 3.48 2.54 3.53 3.98 4.09 4 3.08 2.74 ...
 $ copper  : int  156 54 210 64 143 50 52 52 79 140 ...
 $ alk.phos: num  1718 7395 516 6122 671 ...
 $ ast     : num  137.9 113.5 96.1 60.6 113.2 ...
 $ trig    : int  172 88 55 92 72 63 213 189 88 143 ...
 $ platelet: int  190 221 151 183 136 NA 204 373 251 302 ...
 $ protime : num  12.2 10.6 12 10.3 10.9 11 9.7 11 11 11.5 ...
 $ stage   : int  4 3 4 4 3 3 3 3 2 4 ...
```

```
myLinearModel <- lm(chol ~ bili, data = pbc)
```

```
myLinearModel
```

Call:

```
lm(formula = chol ~ bili, data = pbc)
```


Coefficients:

```
(Intercept)      bili
      303.20      20.24
```

```
summary(myLinearModel)
```

Call:

```
lm(formula = chol ~ bili, data = pbc)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-565.39  -89.90  -35.36   44.92 1285.33
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  303.204      15.601   19.435 < 2e-16 ***
bili         20.240       2.785    7.267 3.63e-12 ***
---

```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 213.2 on 282 degrees of freedom

(134 observations deleted due to missingness)

Multiple R-squared: 0.1577, Adjusted R-squared: 0.1547

F-statistic: 52.8 on 1 and 282 DF, p-value: 3.628e-12

```
broom::tidy(myLinearModel)
```

A tibble: 2 x 5

```
  term          estimate std.error statistic  p.value
<chr>         <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept)   303.      15.6      19.4 5.65e-54
2 bili          20.2      2.79      7.27 3.63e-12
```

```
broom::glance(myLinearModel)
```

A tibble: 1 x 11

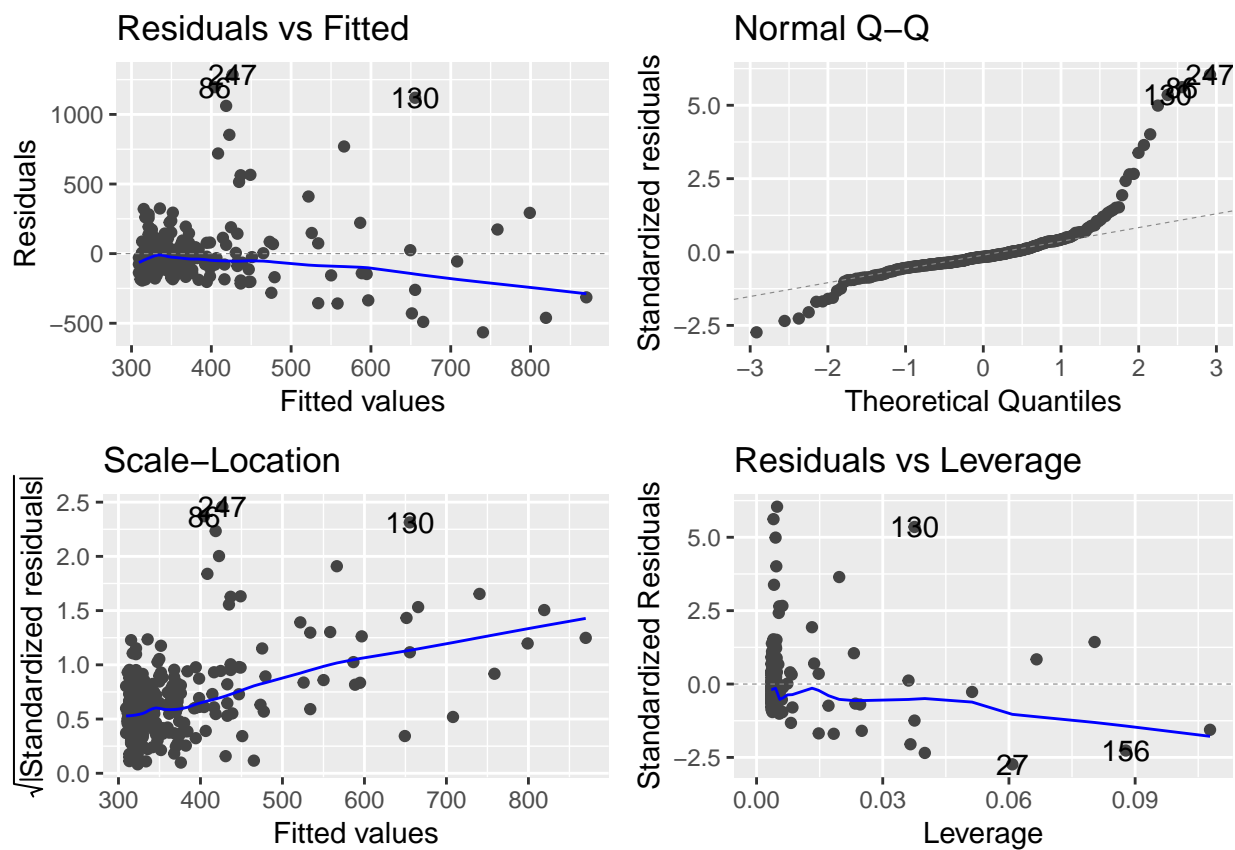
```
  r.squared adj.r.squared sigma statistic  p.value    df logLik  AIC    BIC
    <dbl>         <dbl> <dbl>     <dbl>    <dbl> <int>  <dbl> <dbl> <dbl>
1   0.158         0.155  213.      52.8 3.63e-12     2 -1925. 3856. 3867.
```

... with 2 more variables: deviance <dbl>, df.residual <int>


```
## # install.packages('ggfortify')
```

```
## library(ggfortify)
```


```
## autoplot(myLinearModel)
```



```
## ggplot(pbc, aes(x = bili))+geom_density()
```



```
## ggplot(pbc, aes(x = log(bili)))+geom_density()
```



```
myLinearModel2 <- lm(chol~log(bili), data = pbc)
summary(myLinearModel2)
```

Call:

```
lm(formula = chol ~ log(bili), data = pbc)
```

Residuals:

Min	1Q	Median	3Q	Max
-440.07	-94.35	-21.07	42.67	1221.86

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	311.48	14.28	21.816	< 2e-16 ***
log(bili)	98.80	12.07	8.186	9.42e-15 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

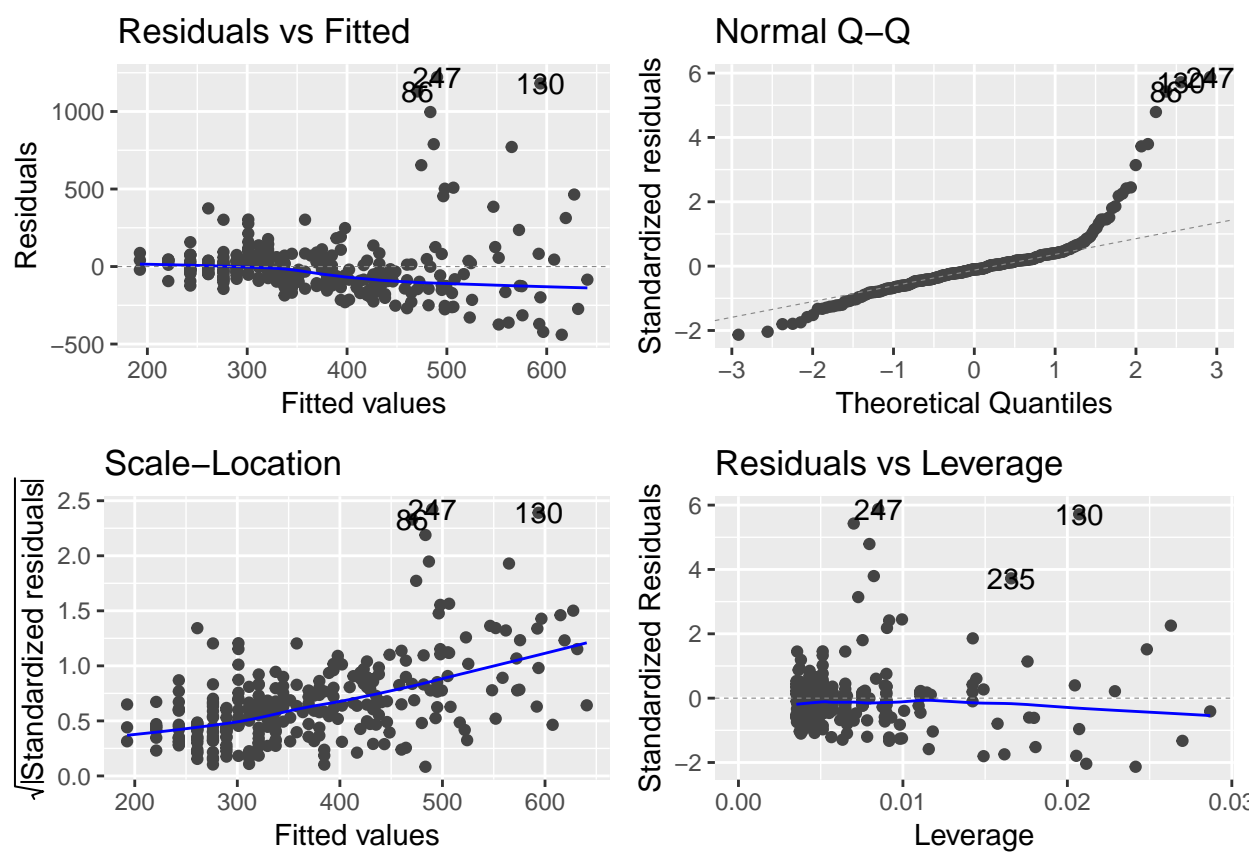
Residual standard error: 208.9 on 282 degrees of freedom

(134 observations deleted due to missingness)


Multiple R-squared: 0.192, Adjusted R-squared: 0.1891

F-statistic: 67.01 on 1 and 282 DF, p-value: 9.416e-15

```
autoplot(myLinearModel2)
```




```
autoplot(myLinearModel2, which=1)
```



```
d <- broom::augment(myLinearModel2)
d
```

```
# A tibble: 284 x 10
  .rownames chol log.bili. .fitted .se.fit .resid .hat .sigma .cooksd
  <chr>      <int>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 1         261      2.67     576.   28.1 -315.  0.0181  208. 2.13e-2
2 2         302     0.0953    321.   13.7 -18.9  0.00433  209. 1.79e-5
3 3         176     0.336     345.   12.8 -169.  0.00373  209. 1.23e-3
4 4         244     0.588     370.   12.4 -126.  0.00352  209. 6.41e-4
5 5         279     1.22     432.   14.6 -153.  0.00487  209. 1.33e-3
6 6         248    -0.223     289.   15.8 -41.4  0.00571  209. 1.14e-4
7 7         322      0       311.   14.3  10.5  0.00467  209. 5.98e-6
8 8         280    -1.20     193.   24.9  87.5  0.0142   209. 1.28e-3
9 9         562     1.16     426.   14.2  136.  0.00463  209. 9.84e-4
10 10        200     2.53     562.   26.6 -362.  0.0162   208. 2.51e-2
# ... with 274 more rows, and 1 more variable: .std.resid <dbl>
```

```
ggplot(d, aes(x = .fitted, y = .resid))+geom_point()+ geom_smooth(se=F)+
  labs(x = 'Fitted values', y = 'Residual values')
```



```
head(predict(myLinearModel2, newdata = pbc))
```

```
1      2      3      4      5      6
575.6925 320.9006 344.7277 369.5578 432.3941 289.4371
```

```
myLM3 <- lm(chol ~ log(bili) + sex, data = pbc)
broom::tidy(myLM3)
```

```
# A tibble: 3 x 5
```

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	283.	36.6	7.71	2.14e-13
2 log(bili)	99.6	12.1	8.22	7.37e-15
3 sexf	32.5	37.8	0.858	3.92e- 1

```
myLR <- glm(spiders ~ albumin + bili + chol, data = pbc, family = binomial)
myLR
```

```
Call: glm(formula = spiders ~ albumin + bili + chol, family = binomial,
  data = pbc)
```

```
Coefficients:
```

```
(Intercept)      albumin          bili          chol
```

```
2.3326484   -0.9954927    0.0995915   -0.0003176
```

```
Degrees of Freedom: 283 Total (i.e. Null); 280 Residual
(134 observations deleted due to missingness)
```

```
Null Deviance:      341.4
```

```
Residual Deviance: 315.2    AIC: 323.2
```

```
broom::tidy(myLR)
```

```
# A tibble: 4 x 5
```

	term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>
1	(Intercept)	2.33	1.30	1.80	0.0717
2	albumin	-0.995	0.362	-2.75	0.00595
3	bili	0.0996	0.0344	2.89	0.00381
4	chol	-0.000318	0.000615	-0.517	0.605

```
broom::glance(myLR)
```

```
# A tibble: 1 x 7
```

	null.deviance <dbl>	df.null <int>	logLik <dbl>	AIC <dbl>	BIC <dbl>	deviance <dbl>	df.residual <int>
1	341.	283	-158.	323.	338.	315.	280

```
head(predict(myLR))
```

	1	2	3	4	5	6
	1.10554163	-1.77506554	-1.04814132	-0.09414055	-0.93144911	-1.62851203

```
head(predict(myLR, type='response'))
```

	1	2	3	4	5	6
	0.7512970	0.1449135	0.2595822	0.4764822	0.2826308	0.1640343

3.1 Model selection

```
# install.packages('leaps')
```

```
library(leaps)
```

```
mtcars1 <- mtcars %>% mutate_at(vars(cyl, vs:carb), as.factor)
```

```
all_subsets <- regsubsets(mpg~., data = mtcars1)
```

```
all_subsets
```

```
Subset selection object
```

```
Call: regsubsets.formula(mpg ~ ., data = mtcars1)
```

```
16 Variables (and intercept)
```

```
Forced in Forced out
```

```
cyl6      FALSE      FALSE
```

```
cyl8      FALSE      FALSE
```

```

disp      FALSE      FALSE
hp        FALSE      FALSE
drat      FALSE      FALSE
wt        FALSE      FALSE
qsec      FALSE      FALSE
vs1       FALSE      FALSE
am1       FALSE      FALSE
gear4     FALSE      FALSE
gear5     FALSE      FALSE
carb2     FALSE      FALSE
carb3     FALSE      FALSE
carb4     FALSE      FALSE
carb6     FALSE      FALSE
carb8     FALSE      FALSE

```

1 subsets of each size up to 8

Selection Algorithm: exhaustive

```

ind <- which.max(summary(all_subsets)$adjr2)
summary(all_subsets)$which[ind,]

```

(Intercept)	cyl6	cyl8	disp	hp	drat
TRUE	TRUE	FALSE	FALSE	TRUE	FALSE
wt	qsec	vs1	am1	gear4	gear5
TRUE	FALSE	TRUE	TRUE	FALSE	FALSE
carb2	carb3	carb4	carb6	carb8	
FALSE	FALSE	FALSE	FALSE	FALSE	

3.2 Many models

```

mtcars <- as_tibble(mtcars)
mtcars %>% select(mpg, disp:qsec)

```

A tibble: 32 x 6

	mpg	disp	hp	drat	wt	qsec
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	21	160	110	3.9	2.62	16.5
2	21	160	110	3.9	2.88	17.0
3	22.8	108	93	3.85	2.32	18.6
4	21.4	258	110	3.08	3.22	19.4
5	18.7	360	175	3.15	3.44	17.0
6	18.1	225	105	2.76	3.46	20.2
7	14.3	360	245	3.21	3.57	15.8
8	24.4	147.	62	3.69	3.19	20
9	22.8	141.	95	3.92	3.15	22.9
10	19.2	168.	123	3.92	3.44	18.3

... with 22 more rows


```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg)
```

```
# A tibble: 160 x 3
   mpg variable value
  <dbl> <chr>   <dbl>
1  21    disp    160
2  21    disp    160
3 22.8   disp    108
4 21.4   disp    258
5 18.7   disp    360
6 18.1   disp    225
7 14.3   disp    360
8 24.4   disp    147.
9 22.8   disp    141.
10 19.2   disp    168.
# ... with 150 more rows
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  group_by(variable) %>%
  lm(mpg~value, data=.)
```

Call:

```
lm(formula = mpg ~ value, data = .)
```

Coefficients:

```
(Intercept)      value
 21.28328      -0.01483
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable)
```

```
# A tibble: 5 x 2
  variable data
  <chr>    <list>
1 disp    <tibble [32 x 2]>
2 hp      <tibble [32 x 2]>
3 drat    <tibble [32 x 2]>
4 wt      <tibble [32 x 2]>
5 qsec    <tibble [32 x 2]>
```

```
bl <- mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable)
bl$data[[1]]
```

```
# A tibble: 32 x 2
```

```
  mpg value
<dbl> <dbl>
1  21    160
2  21    160
3  22.8  108
4  21.4  258
5  18.7  360
6  18.1  225
7  14.3  360
8  24.4  147.
9  22.8  141.
10 19.2  168.
```

```
# ... with 22 more rows
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable) %>%
  mutate(models = map(data, ~lm(mpg~value, data=.)))
```

```
# A tibble: 5 x 3
```

```
  variable data          models
  <chr>    <list>         <list>
1 disp    <tibble [32 x 2]> <S3: lm>
2 hp      <tibble [32 x 2]> <S3: lm>
3 drat    <tibble [32 x 2]> <S3: lm>
4 wt      <tibble [32 x 2]> <S3: lm>
5 qsec    <tibble [32 x 2]> <S3: lm>
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable) %>%
  mutate(models = map(data, ~lm(mpg~value, data=.) ),
         outputs = map(models, ~tidy(.)))
```

```
# A tibble: 5 x 4
```

```
  variable data          models  outputs
  <chr>    <list>         <list>  <list>
1 disp    <tibble [32 x 2]> <S3: lm> <tibble [2 x 5]>
2 hp      <tibble [32 x 2]> <S3: lm> <tibble [2 x 5]>
3 drat    <tibble [32 x 2]> <S3: lm> <tibble [2 x 5]>
4 wt      <tibble [32 x 2]> <S3: lm> <tibble [2 x 5]>
5 qsec    <tibble [32 x 2]> <S3: lm> <tibble [2 x 5]>
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable) %>%
  mutate(models = map(data, ~lm(mpg~value, data=.) ),
         outputs = map(models, ~tidy(.))) %>%
```

```
select(variable, outputs)
```

```
# A tibble: 5 x 2
  variable outputs
  <chr>      <list>
1 disp      <tibble [2 x 5]>
2 hp        <tibble [2 x 5]>
3 drat      <tibble [2 x 5]>
4 wt        <tibble [2 x 5]>
5 qsec      <tibble [2 x 5]>
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable) %>%
  mutate(models = map(data, ~lm(mpg~value, data=)),
        outputs = map(models, ~tidy(.))) %>%
  select(variable, outputs) %>%
  unnest()
```

```
# A tibble: 10 x 6
  variable term      estimate std.error statistic p.value
  <chr>      <chr>      <dbl>      <dbl>      <dbl>      <dbl>
1 disp      (Intercept)  29.6        1.23        24.1  3.58e-21
2 disp      value      -0.0412     0.00471    -8.75  9.38e-10
3 hp        (Intercept)  30.1        1.63        18.4  6.64e-18
4 hp        value      -0.0682     0.0101     -6.74  1.79e- 7
5 drat      (Intercept) -7.52        5.48        -1.37  1.80e- 1
6 drat      value       7.68        1.51         5.10  1.78e- 5
7 wt        (Intercept)  37.3        1.88        19.9  8.24e-19
8 wt        value      -5.34        0.559      -9.56  1.29e-10
9 qsec      (Intercept) -5.11       10.0        -0.510 6.14e- 1
10 qsec     value       1.41        0.559       2.53  1.71e- 2
```

```
mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable) %>%
  mutate(models = map(data, ~lm(mpg~value, data=)),
        outputs = map(models, ~tidy(.))) %>%
  select(variable, outputs) %>%
  unnest() %>%
  filter(term=='value')
```

```
# A tibble: 5 x 6
  variable term estimate std.error statistic p.value
  <chr>      <chr>      <dbl>      <dbl>      <dbl>      <dbl>
1 disp      value -0.0412  0.00471    -8.75  9.38e-10
2 hp        value -0.0682  0.0101     -6.74  1.79e- 7
3 drat      value  7.68     1.51         5.10  1.78e- 5
```

```

4 wt      value  -5.34      0.559      -9.56 1.29e-10
5 qsec    value   1.41      0.559       2.53 1.71e- 2

```

```

mtcars %>% select(mpg, disp:qsec) %>%
  gather(variable, value, -mpg) %>%
  nest(-variable) %>%
  mutate(models = map(data, ~lm(mpg~value, data=.)),
         outputs = map(models, ~tidy(.))) %>%
  select(variable, outputs) %>%
  unnest() %>%
  filter(term=='value') %>%
  mutate_if(is.numeric, funs(round(., 3)))

```

```

# A tibble: 5 x 6
  variable term estimate std.error statistic p.value
  <chr>    <chr>    <dbl>    <dbl>    <dbl>    <dbl>
1 disp    value    -0.041    0.005    -8.75     0
2 hp      value    -0.068    0.01     -6.74     0
3 drat    value     7.68     1.51     5.10     0
4 wt      value    -5.34     0.559    -9.56     0
5 qsec    value     1.41     0.559     2.52    0.017

```

Chapter 4

Predictive modeling

```
library(tidyverse)
library(caret)
data(diamonds)
set.seed(12356)
diamonds_train <- diamonds %>% sample_frac(size = 0.8) # 80%
diamonds_test <- anti_join(diamonds, diamonds_train)
(nrow(diamonds) == nrow(diamonds_train) + nrow(diamonds_test))
```

```
[1] FALSE
```